

The Problem Statement:

To build an application to classify the patients to be healthy or suffering from cardiovascular disease based on the given attributes.

Features:

1. Age | Objective Feature | age | int (days)
2. Height | Objective Feature | height | int (cm) |
3. Weight | Objective Feature | weight | float (kg) |
4. Gender | Objective Feature | gender | categorical code |
5. Systolic blood pressure | Examination Feature | ap_hi | int |
6. Diastolic blood pressure | Examination Feature | ap_lo | int |
7. Cholesterol | Examination Feature | cholesterol | 1: normal, 2: above normal, 3: well above normal |
8. Glucose | Examination Feature | gluc | 1: normal, 2: above normal, 3: well above normal |
9. Smoking | Subjective Feature | smoke | binary |
10. Alcohol intake | Subjective Feature | alco | binary |
11. Physical activity | Subjective Feature | active | binary |
12. Presence or absence of cardiovascular disease | Target Variable | cardio | binary |

All of the dataset values were collected at the moment of medical examination.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier
```

```
In [2]: # Load the data
cardio = pd.read_csv('cardio_train.csv', delimiter=';')
cardio.head(3)
```

```
Out[2]:
```

	id	age	gender	height	weight	ap_hi	ap_lo	cholesterol	gluc	smoke	alco	active	cardio
0	988	22469	1	155	69.0	130	80	2	2	0	0	1	0
1	989	14648	1	163	71.0	110	70	1	1	0	0	1	1
2	990	21901	1	165	70.0	120	80	1	1	0	0	1	0

```
In [3]: cardio.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 69301 entries, 0 to 69300
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               69301 non-null  int64
1   age              69301 non-null  int64
2   gender           69301 non-null  int64
3   height           69301 non-null  int64
4   weight           69301 non-null  float64
5   ap_hi            69301 non-null  int64
6   ap_lo            69301 non-null  int64
7   cholesterol      69301 non-null  int64
8   gluc             69301 non-null  int64
9   smoke            69301 non-null  int64
10  alco             69301 non-null  int64
11  active           69301 non-null  int64
12  cardio           69301 non-null  int64
dtypes: float64(1), int64(12)
memory usage: 6.9 MB
```

Age, height, weight, ap_hi, ap_lo are continues data

cholesterol, gluc, smoke, alco, active, cardio are categorical data

id is not required column

```
In [4]: cardio.drop(columns=['id'], inplace=True)
```

```
In [5]: cardio.describe().T
```

Out[5]:

	count	mean	std	min	25%	50%	75%	max
age	69301.0	19468.786280	2467.261818	10798.0	17664.0	19704.0	21326.0	23713.0
gender	69301.0	1.349519	0.476821	1.0	1.0	1.0	2.0	2.0
height	69301.0	164.362217	8.205337	55.0	159.0	165.0	170.0	250.0
weight	69301.0	74.203027	14.383469	10.0	65.0	72.0	82.0	200.0
ap_hi	69301.0	128.829584	154.775805	-150.0	120.0	120.0	140.0	16020.0
ap_lo	69301.0	96.650092	189.096240	-70.0	80.0	80.0	90.0	11000.0
cholesterol	69301.0	1.366806	0.680270	1.0	1.0	1.0	2.0	3.0
gluc	69301.0	1.226447	0.572246	1.0	1.0	1.0	1.0	3.0
smoke	69301.0	0.088051	0.283371	0.0	0.0	0.0	0.0	1.0
alco	69301.0	0.053881	0.225784	0.0	0.0	0.0	0.0	1.0
active	69301.0	0.803986	0.396982	0.0	1.0	1.0	1.0	1.0
cardio	69301.0	0.499589	0.500003	0.0	0.0	0.0	1.0	1.0

as looking into count column we can see there is no null values present in any column

lets verify

```
In [6]: cardio.isna().sum()
```

```
Out[6]: age          0
gender        0
height        0
weight        0
ap_hi         0
ap_lo         0
cholesterol    0
gluc          0
smoke         0
alco          0
active        0
cardio        0
dtype: int64
```


Age in days is not easy to understand

Lets make new column called year where age in days is converted into age in year

```
In [7]: cardio['age'] = (cardio['age'] / 365).round().astype('int')
```

Lets Plot a boxplot to check age(age in year)column

```
In [8]: plt.figure(figsize=(4,7), facecolor='DarkOrange')  
sns.boxplot(cardio.age, color='Brown');
```



Describe function

Difference between mean and standard deviation shows that there so many variances in data,

Lets only consider years age above quantile 0.015, they are only few below age 40

Lets consider only 1.5% to 97.5% of weight and height and 1.5% above age

```
In [9]: qua_low = cardio.quantile(0.015)
qua_high = cardio.quantile(0.975)

c1 = cardio[(cardio.height > qua_high.height) |
             (cardio.height < qua_low.height) |
             (cardio.weight > qua_high.weight) |
             (cardio.weight < qua_low.weight) |
             (cardio.age < qua_low.age)].index

cardio.drop(c1, inplace=True)
```

ap_hi and ap_lo represents blood pressure,

difference between ap_hi and ap_lo show the blood pressure,

while ap_hi and ap_lo cant be negative

```
In [10]: cardio.ap_hi = cardio.ap_hi.abs()
cardio.ap_lo = cardio.ap_lo.abs()
```

There are some outlier data present in both columns

Lets consider only 1.5% to 97.5% of ap_hi and ap_lo


```
In [11]: cardio.drop(cardio[(cardio.ap_hi > qua_high.ap_hi) |
                             (cardio.ap_hi < qua_low.ap_hi) |
                             (cardio.ap_lo > qua_high.ap_lo) |
                             (cardio.ap_lo < qua_low.ap_lo)].index, inplace=True)
```

```
In [12]: cardio_cont= cardio[['age', 'height', 'weight', 'ap_hi', 'ap_lo']]

plt.figure(figsize=(15,10), facecolor='BurlyWood')
plotnumber =1

for column in cardio_cont:
    if plotnumber<=5:
        ax=plt.subplot(1,5,plotnumber)
        sns.boxplot(cardio_cont[column], color='Coral')
        plt.xlabel(column.upper(), fontsize = 20)

        plotnumber+=1
plt.tight_layout()
plt.show()
```



In [13]: `cardio.describe().T`

Out[13]:


	count	mean	std	min	25%	50%	75%	max
age	61902.0	53.400827	6.686393	40.0	49.0	54.0	58.0	65.0
gender	61902.0	1.341104	0.474084	1.0	1.0	1.0	2.0	2.0
height	61902.0	164.313382	6.990518	148.0	159.0	165.0	169.0	180.0
weight	61902.0	73.212112	11.982543	50.0	65.0	71.0	80.0	108.0
ap_hi	61902.0	125.676424	14.823129	90.0	120.0	120.0	140.0	170.0
ap_lo	61902.0	80.899890	8.581824	60.0	80.0	80.0	90.0	100.0
cholesterol	61902.0	1.354205	0.671738	1.0	1.0	1.0	1.0	3.0
gluc	61902.0	1.220300	0.566947	1.0	1.0	1.0	1.0	3.0
smoke	61902.0	0.084876	0.278699	0.0	0.0	0.0	0.0	1.0
alco	61902.0	0.051501	0.221019	0.0	0.0	0.0	0.0	1.0
active	61902.0	0.804174	0.396838	0.0	1.0	1.0	1.0	1.0
cardio	61902.0	0.487916	0.499858	0.0	0.0	0.0	1.0	1.0

Plotting continues data vs cardio data

```
In [14]: cont= cardio[['age','height','weight','ap_hi','ap_lo']]

plt.figure(figsize=(25,10), facecolor='Tomato')
plotnumber = 1

for columns in cont:
    if plotnumber <=5:
        ax=plt.subplot(1,5,plotnumber)
        #sns.boxenplot(cont[columns])
        sns.violinplot(data=cardio, y=cont[columns], x="cardio")
        plt.ylabel(columns.upper(), fontsize = 15)
        plt.xlabel('Cardio', fontsize= 15)
        plotnumber+=1
```


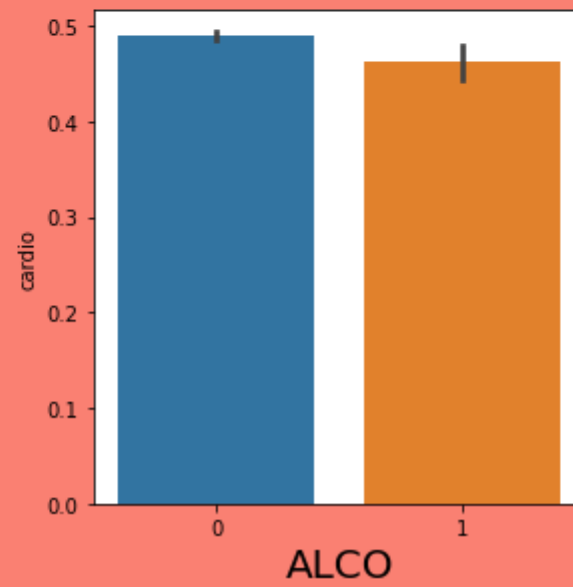
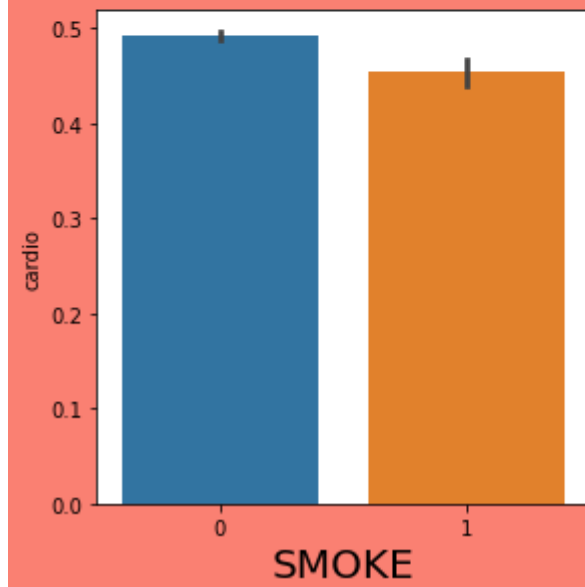
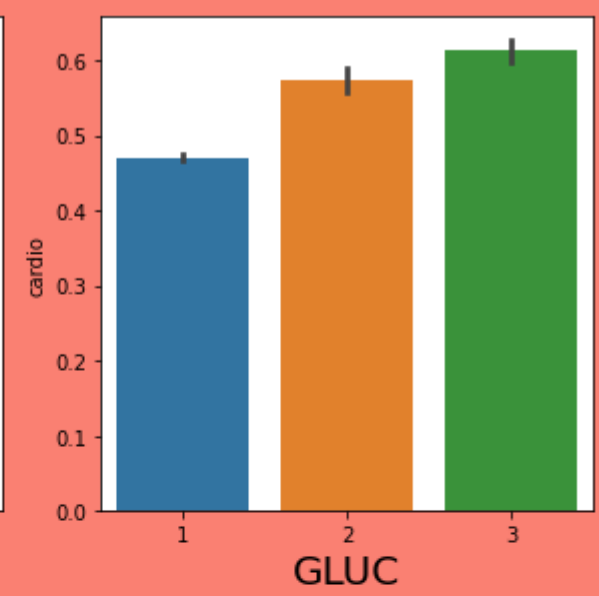




Plotting categorical data vs cardio data

```
In [15]: cate= cardio[['gender','cholesterol','gluc','smoke','alco','active']]

plt.figure(figsize=(15,10), facecolor='Salmon')
plotnumber = 1




for columns in cate:
    if plotnumber <=6:
        ax=plt.subplot(2,3,plotnumber)
        sns.barplot(x=cate[columns], y=cardio.cardio, data= cardio)
        plt.xlabel(columns.upper(), fontsize = 20)
        plotnumber+=1
```

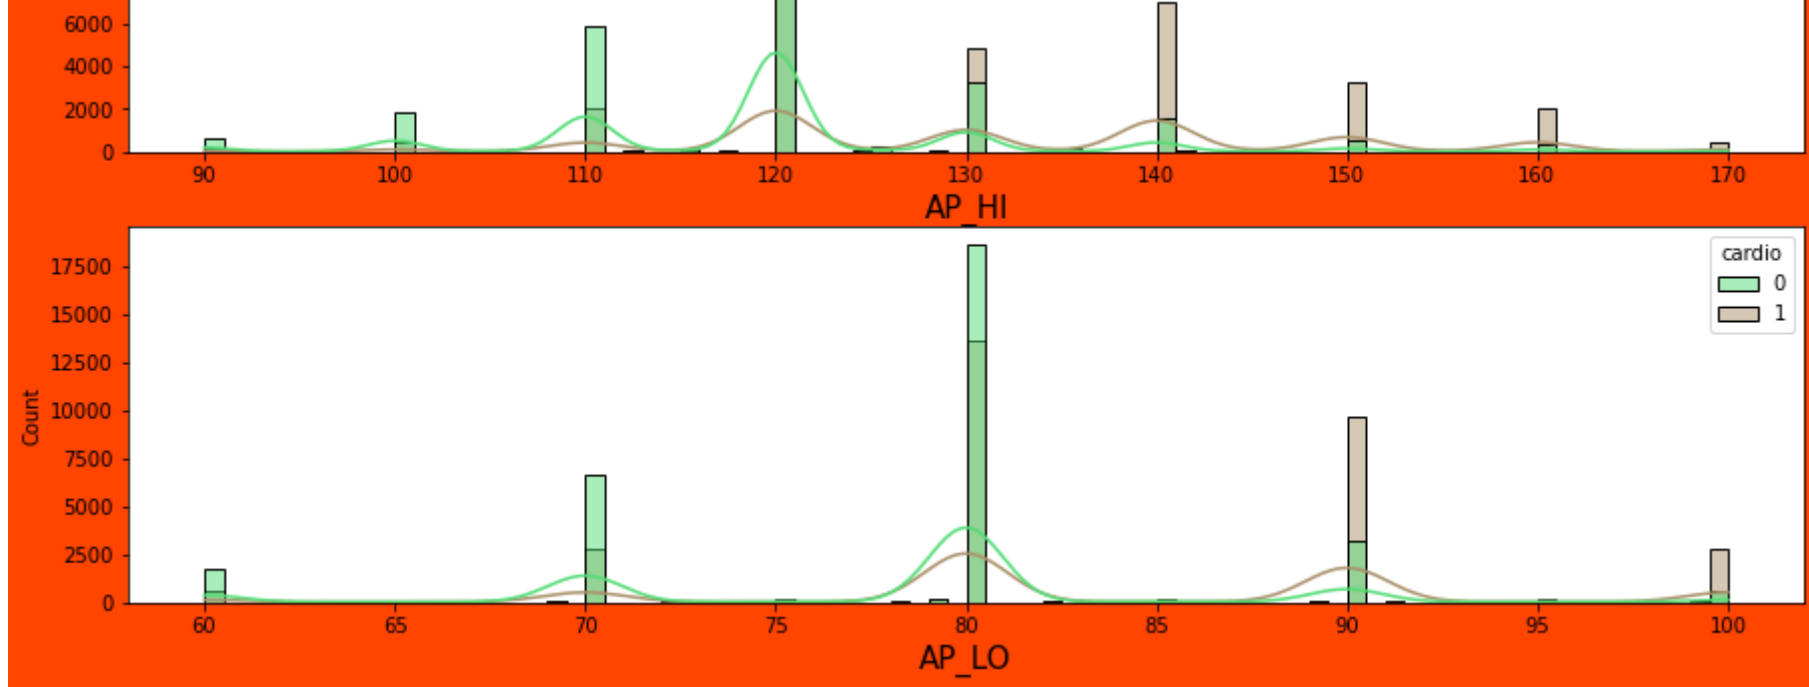



```
In [16]: cont= cardio[['age','height','weight','ap_hi','ap_lo']]

plt.figure(figsize=(15,20), facecolor='OrangeRed')
plotnumber = 1


for columns in cont:
    if plotnumber <=5:
        ax=plt.subplot(5,1,plotnumber)
        sns.histplot(x= cont[columns], hue= cardio.cardio,
                     data=cardio, palette='terrain', kde=True)
        plt.xlabel(columns.upper(), fontsize = 15)
        plotnumber+=1
```





Plot heatmap with correlation of cardio

```
In [17]: plt.figure(figsize=(15,15), facecolor='SandyBrown')
sns.heatmap(cardio.corr(), annot=True, cmap='coolwarm',linewidths=0.05);
```





Data looks good,

Lets begin Model Training

Classification

Logistic Regression model to predict the cardio

```
In [18]: # Separating feature and output/result
X = cardio.drop(columns=['cardio'])
y = cardio.cardio

# Splting train and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=100)

# Initializing Logistic regression
lgr = LogisticRegression()

# Fitting Logistic modle
lgr.fit(X_train, y_train)

# Predict the model
y_pred = lgr.predict(X_test)

# Model Report/results
print('Accuracy Score', (accuracy_score(y_pred,y_test))*100)
print('\n Confusion Matrix \n', confusion_matrix(y_pred, y_test))
print('\n Classification Report \n', classification_report(y_pred,y_test))
```

Accuracy Score 70.99379684673042

Confusion Matrix

[[6100 2617]

[1872 4887]]

Classification Report

	precision	recall	f1-score	support
0	0.77	0.70	0.73	8717
1	0.65	0.72	0.69	6759
accuracy			0.71	15476
macro avg	0.71	0.71	0.71	15476
weighted avg	0.72	0.71	0.71	15476

Accuracy of the model is 71%

Lets try Standard Scaler to check accuracy

Standard Scaler

```
In [19]: # Initialize Scaler Model
scaler = StandardScaler()

# Apply scaler model
X_scaled = scaler.fit_transform(X)

# Splitting data into train and test data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.25, random_state=100)

# Fitting Logistic modle
lgr.fit(X_train, y_train)

# Predict the model
y_pred = lgr.predict(X_test)

# Model Report/results
print('Accuracy Score', (accuracy_score(y_pred,y_test))*100)
print('\n Confusion Matrix \n', confusion_matrix(y_pred, y_test))
print('\n Classification Report \n', classification_report(y_pred,y_test))
```

Accuracy Score 71.97596278108038

Confusion Matrix

```
[[6249 2614]
 [1723 4890]]
```

Classification Report

	precision	recall	f1-score	support
0	0.78	0.71	0.74	8863
1	0.65	0.74	0.69	6613
accuracy			0.72	15476
macro avg	0.72	0.72	0.72	15476
weighted avg	0.73	0.72	0.72	15476

Accuracy is 72% using standard scaler the result variance is not that much affected

Lets try Decision Tree and see

Decision Tree

```
In [20]: # Splitting data into train and test data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.25, random_state=250)

# Initializing Decision Tree
Class_tree = DecisionTreeClassifier(criterion= 'gini', max_depth= 3)

# Applying Decision Tree model to data
Class_tree.fit(X_train,y_train)

# Predict the model
y_pred = Class_tree.predict(X_test)

# Model Report/results
print('Accuracy Score', (accuracy_score(y_pred,y_test))*100)
print('\n Confusion Matrix \n', confusion_matrix(y_pred, y_test))
print('\n Classification Report \n', classification_report(y_pred,y_test))
```


Accuracy Score 72.08581028689584

Confusion Matrix

[[6118 2632]

[1688 5038]]

Classification Report

	precision	recall	f1-score	support
0	0.78	0.70	0.74	8750
1	0.66	0.75	0.70	6726
accuracy			0.72	15476
macro avg	0.72	0.72	0.72	15476
weighted avg	0.73	0.72	0.72	15476

accuracy is still same.

lets do hyperparameter- tuning to see which tree gives best results

Hyperparameter Tuning

```
In [21]: # parameter's
param = {
    'criterion': ['gini','entropy',"log_loss"],
    'max_depth' : range(6,15),
    'min_samples_leaf' : range(11,16),
    'max_features' : [ 'sqrt', 'log2' , None],
    'splitter' : ['best','random']
}

# Initializing Grid Search CV
grid_search = GridSearchCV(estimator=Class_tree,
                           param_grid=param, cv=5, n_jobs=-1)

#Applying Grid Search CV
grid_search.fit(X_train, y_train)

#Result
print('Grid Search Best Parameter',grid_search.best_params_)
print('Grid Search Best Score',grid_search.best_score_)
```

```
Grid Search Best Parameter {'criterion': 'gini', 'max_depth': 8, 'max_features': None, 'min_samples_leaf': 15, 'splitter': 'random'}
Grid Search Best Score 0.7281480496925848
```

```
In [22]: # Initializing Decision Tree
Class_tree = DecisionTreeClassifier(criterion= 'entropy', max_depth= 8, max_features = None,
                                   min_samples_leaf= 15, splitter='random')

# Applying Decision Tree model to data
Class_tree.fit(X_train,y_train)

# Predict the model
y_pred = Class_tree.predict(X_test)

# Model Report/results
print('Accuracy Score', (accuracy_score(y_pred,y_test))*100)
print('\n Confusion Matrix \n', confusion_matrix(y_pred, y_test))
print('\n Classification Report \n', classification_report(y_pred,y_test))
```

Accuracy Score 72.40242956836391

Confusion Matrix

```
[[6338 2803]
 [1468 4867]]
```

Classification Report

	precision	recall	f1-score	support
0	0.81	0.69	0.75	9141
1	0.63	0.77	0.70	6335
accuracy			0.72	15476
macro avg	0.72	0.73	0.72	15476
weighted avg	0.74	0.72	0.73	15476

Accuracy Score is 72.4, even after hyperparameter tuning the result is same