

The Problem Statement:

To build an application which predicts the monthly rental of a house based on the given attributes.

Attributes:

- Id: listing id
- url: listing URL
- region: craigslist region
- region_url: region URL
- price: rent per month (Target Column)
- type: housing type
- sqfeet: total square footage
- beds:number of beds
- baths:number of bathrooms
- cats_allowed: cats allowed boolean (1 = yes, 0 = no)
- dogs_allowed: dogs allowed boolean
- smoking_allowed: smoking allowed boolean
- wheelchair_access: has wheelchair access boolean
- electric_vehicle_charge: has electric vehicle charger boolean
- comes_furnished: comes with furniture boolean
- laundry_options: laundry options available
- parking_options: parking options available
- image_url: image URL
- description: description by poster
- lat: latitude
- long: longitude
- state: state of listing

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.tree import DecisionTreeRegressor

import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: house = pd.read_csv("housing_train.csv")
house.head(3)
```

Out[2]:

	id	url	region	region_url	price	type	sqfeet	beds	baths	cats_allowed	...	wheelchair_access	electric_ve
0	7039061606	https://bham.craigslist.org/apa/d/birmingham-h...	birmingham	https://bham.craigslist.org	1195	apartment	1908	3	2.0	1	...	0	
1	7041970863	https://bham.craigslist.org/apa/d/birmingham-w...	birmingham	https://bham.craigslist.org	1120	apartment	1319	3	2.0	1	...	0	
2	7041966914	https://bham.craigslist.org/apa/d/birmingham-g...	birmingham	https://bham.craigslist.org	825	apartment	1133	1	1.5	1	...	0	

3 rows × 22 columns

```
In [3]: house.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 265190 entries, 0 to 265189
Data columns (total 22 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    265190 non-null  int64
1   url                                    265190 non-null  object
2   region                                265190 non-null  object
3   region_url                            265190 non-null  object
4   price                                 265190 non-null  int64
5   type                                  265190 non-null  object
6   sqfeet                                265190 non-null  int64
7   beds                                  265190 non-null  int64
8   baths                                 265190 non-null  float64
9   cats_allowed                          265190 non-null  int64
10  dogs_allowed                          265190 non-null  int64
11  smoking_allowed                       265190 non-null  int64
12  wheelchair_access                     265190 non-null  int64
13  electric_vehicle_charge                265190 non-null  int64
14  comes_furnished                       265190 non-null  int64
15  laundry_options                        210879 non-null  object
16  parking_options                        170055 non-null  object
17  image_url                              265190 non-null  object
18  description                             265188 non-null  object
19  lat                                    263771 non-null  float64
20  long                                    263771 non-null  float64
21  state                                  265189 non-null  object
dtypes: float64(3), int64(10), object(9)
memory usage: 44.5+ MB
```

Drop unwanted columns

```
In [4]: house.drop(columns=['id','url','lat','long','region_url','image_url','description'], inplace=True)
```

Shape of the data

```
In [5]: house.shape
```

Out[5]: (265190, 15)

Lets check the statical values of data

```
In [6]: house.describe().round(2).T
```

Out[6]:

		count	mean	std	min	25%	50%	75%	max
	price	265190.0	12272.85	5376351.72	0.0	817.0	1060.0	1450.0	2.768307e+09
	sqfeet	265190.0	1093.68	23068.88	0.0	752.0	950.0	1156.0	8.388607e+06
	beds	265190.0	1.91	3.69	0.0	1.0	2.0	2.0	1.100000e+03
	baths	265190.0	1.48	0.63	0.0	1.0	1.0	2.0	7.500000e+01
	cats_allowed	265190.0	0.72	0.45	0.0	0.0	1.0	1.0	1.000000e+00
	dogs_allowed	265190.0	0.70	0.46	0.0	0.0	1.0	1.0	1.000000e+00
	smoking_allowed	265190.0	0.73	0.44	0.0	0.0	1.0	1.0	1.000000e+00
	wheelchair_access	265190.0	0.08	0.27	0.0	0.0	0.0	0.0	1.000000e+00
	electric_vehicle_charge	265190.0	0.01	0.12	0.0	0.0	0.0	0.0	1.000000e+00
	comes_furnished	265190.0	0.05	0.22	0.0	0.0	0.0	0.0	1.000000e+00

Checking null values

```
In [7]: house.isna().sum()
```

Out[7]:

region	0
price	0
type	0
sqfeet	0
beds	0
baths	0
cats_allowed	0
dogs_allowed	0
smoking_allowed	0
wheelchair_access	0
electric_vehicle_charge	0
comes_furnished	0
laundry_options	54311
parking_options	95135
state	1

dtype: int64

Remove Null values and Duplicate

```
In [8]: house.dropna(inplace=True)
```

```
In [9]: house.drop_duplicates(inplace=True)
```

Check the unique values to see the data in categorical values and some continuous values

```
In [10]: print('Unique values of beds \n ',house.beds.unique(),'\n',
              'Unique values of baths \n ',house.baths.unique(),'\n',
              'Available laundry options \n ',house.laundry_options.unique(),'\n',
              'Available parking options \n ',house.parking_options.unique())
```

```
Unique values of beds
[ 3  1  2  4  0  5  7  6  8 1000 1100]
Unique values of baths
[ 2.  1.5  1.  2.5  3.5  0.  3.  4.  4.5  5.  6.5  6.  5.5  7.
35. 75.  7.5  8. ]
Available laundry options
['laundry on site' 'w/d hookups' 'laundry in bldg' 'w/d in unit'
'no laundry on site']
Available parking options
['street parking' 'off-street parking' 'carport' 'attached garage'
'detached garage' 'no parking' 'valet parking']
```

Baths values are uneven and lets replace with next values of the uneven values

```
In [11]: house.baths.replace(to_replace=(1.5,2.5,3.5,4.5), value=(2,3,4,5), inplace=True )
```

Describe function

Difference between mean and standard deviation shows that there so many variances in data,

Min and Max Values of Sqfeet, Beds, Baths and Price are to extermes, seems like there are outliers

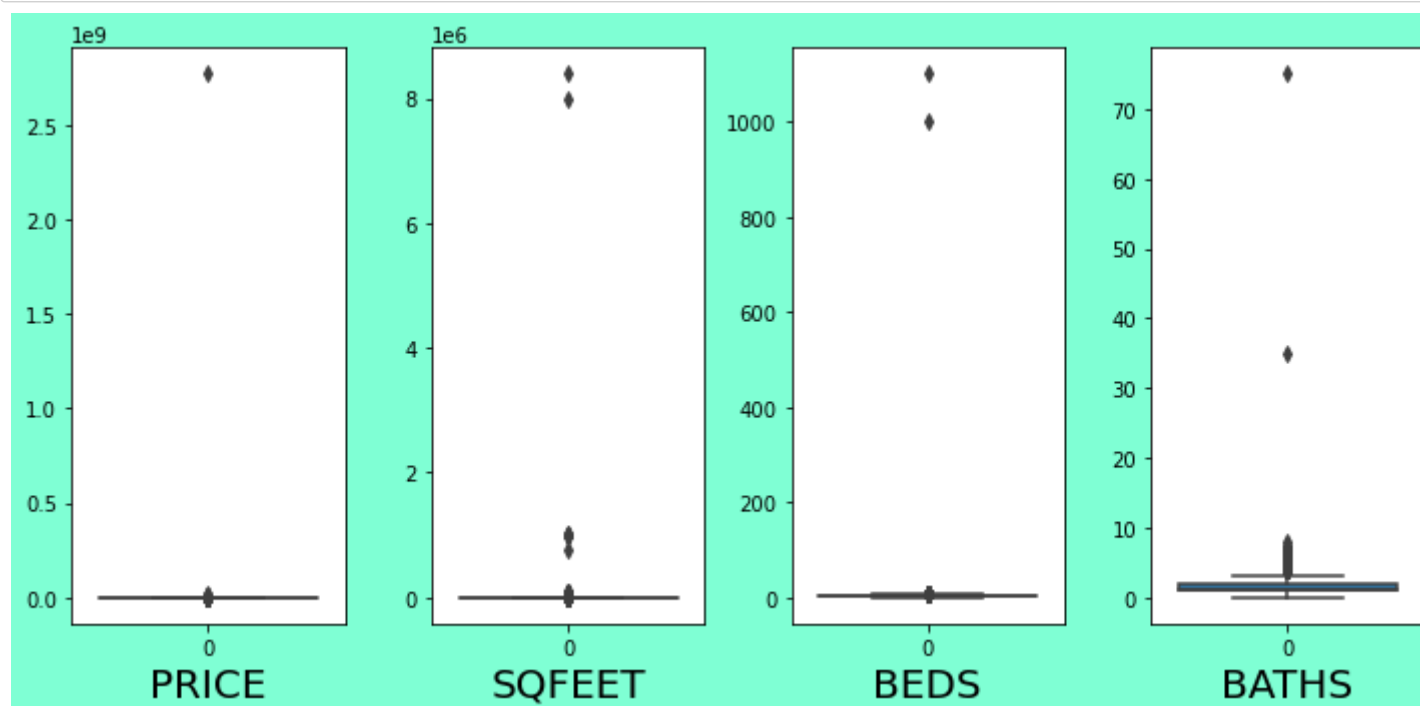
lets check data with box plot

```
In [12]: house_cont= house[['price','sqfeet','beds','baths']]

plt.figure(figsize=(10,5), facecolor='aquamarine')
plotnumber =1

for column in house_cont:
    if plotnumber<=4:
        ax=plt.subplot(1,4,plotnumber)
        sns.boxplot(house_cont[column])
        plt.xlabel(column.upper(), fontsize = 20)

        plotnumber+=1
plt.tight_layout()
plt.show()
```



Bar plot can also shows the outliers present in Price, Sqfeet, Beds and Baths

Lets remove data of price, sqfeet, beds and baths which isn't less than 12.5% and more than 97.5%

```
In [13]: qua_low = house.quantile(0.125)
qua_high = house.quantile(0.975)

h1= house[(house.price>qua_high.price)|
          (house.price<qua_low.price)|
          (house.sqfeet>qua_high.sqfeet)|
          (house.sqfeet<qua_low.sqfeet)|
          (house.beds>qua_high.beds)|
          (house.beds<qua_low.beds)|
          (house.baths>qua_high.baths)|
          (house.baths<qua_low.baths)].index

house.drop(h1, inplace=True)
```

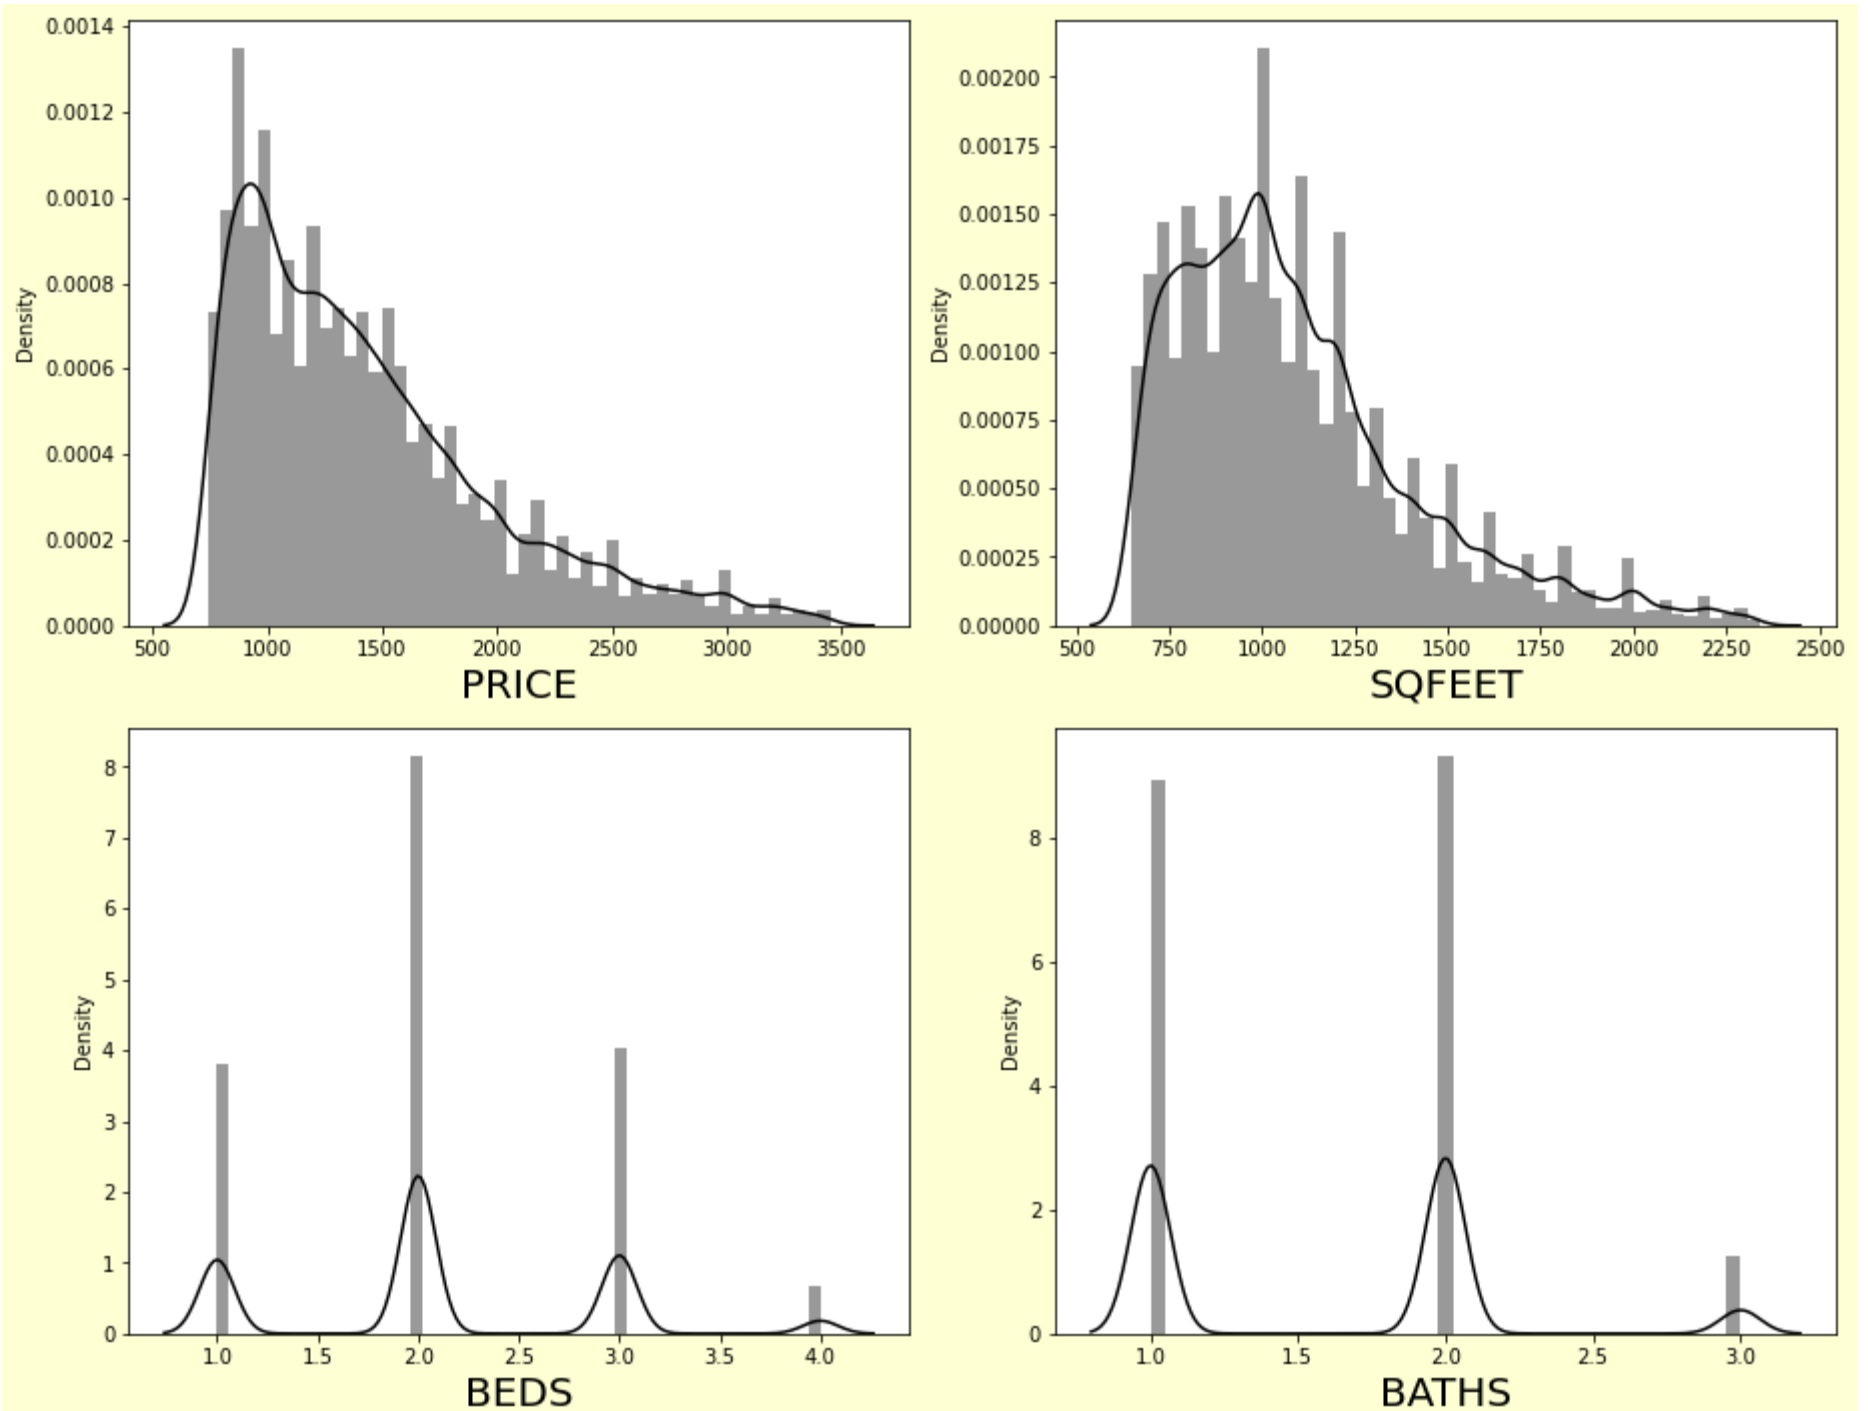
Lets see the Disturbance plot and how values are distributed

```
In [14]: house_cont= house[['price', 'sqfeet', 'beds', 'baths']]

plt.figure(figsize=(13,10), facecolor='xkcd:eggshell')
plotnumber =1

for column in house_cont:
    if plotnumber<=10:
        ax=plt.subplot(2,2,plotnumber)
        sns.distplot(house_cont[column], color='k')
        plt.xlabel(column.upper(), fontsize = 20)

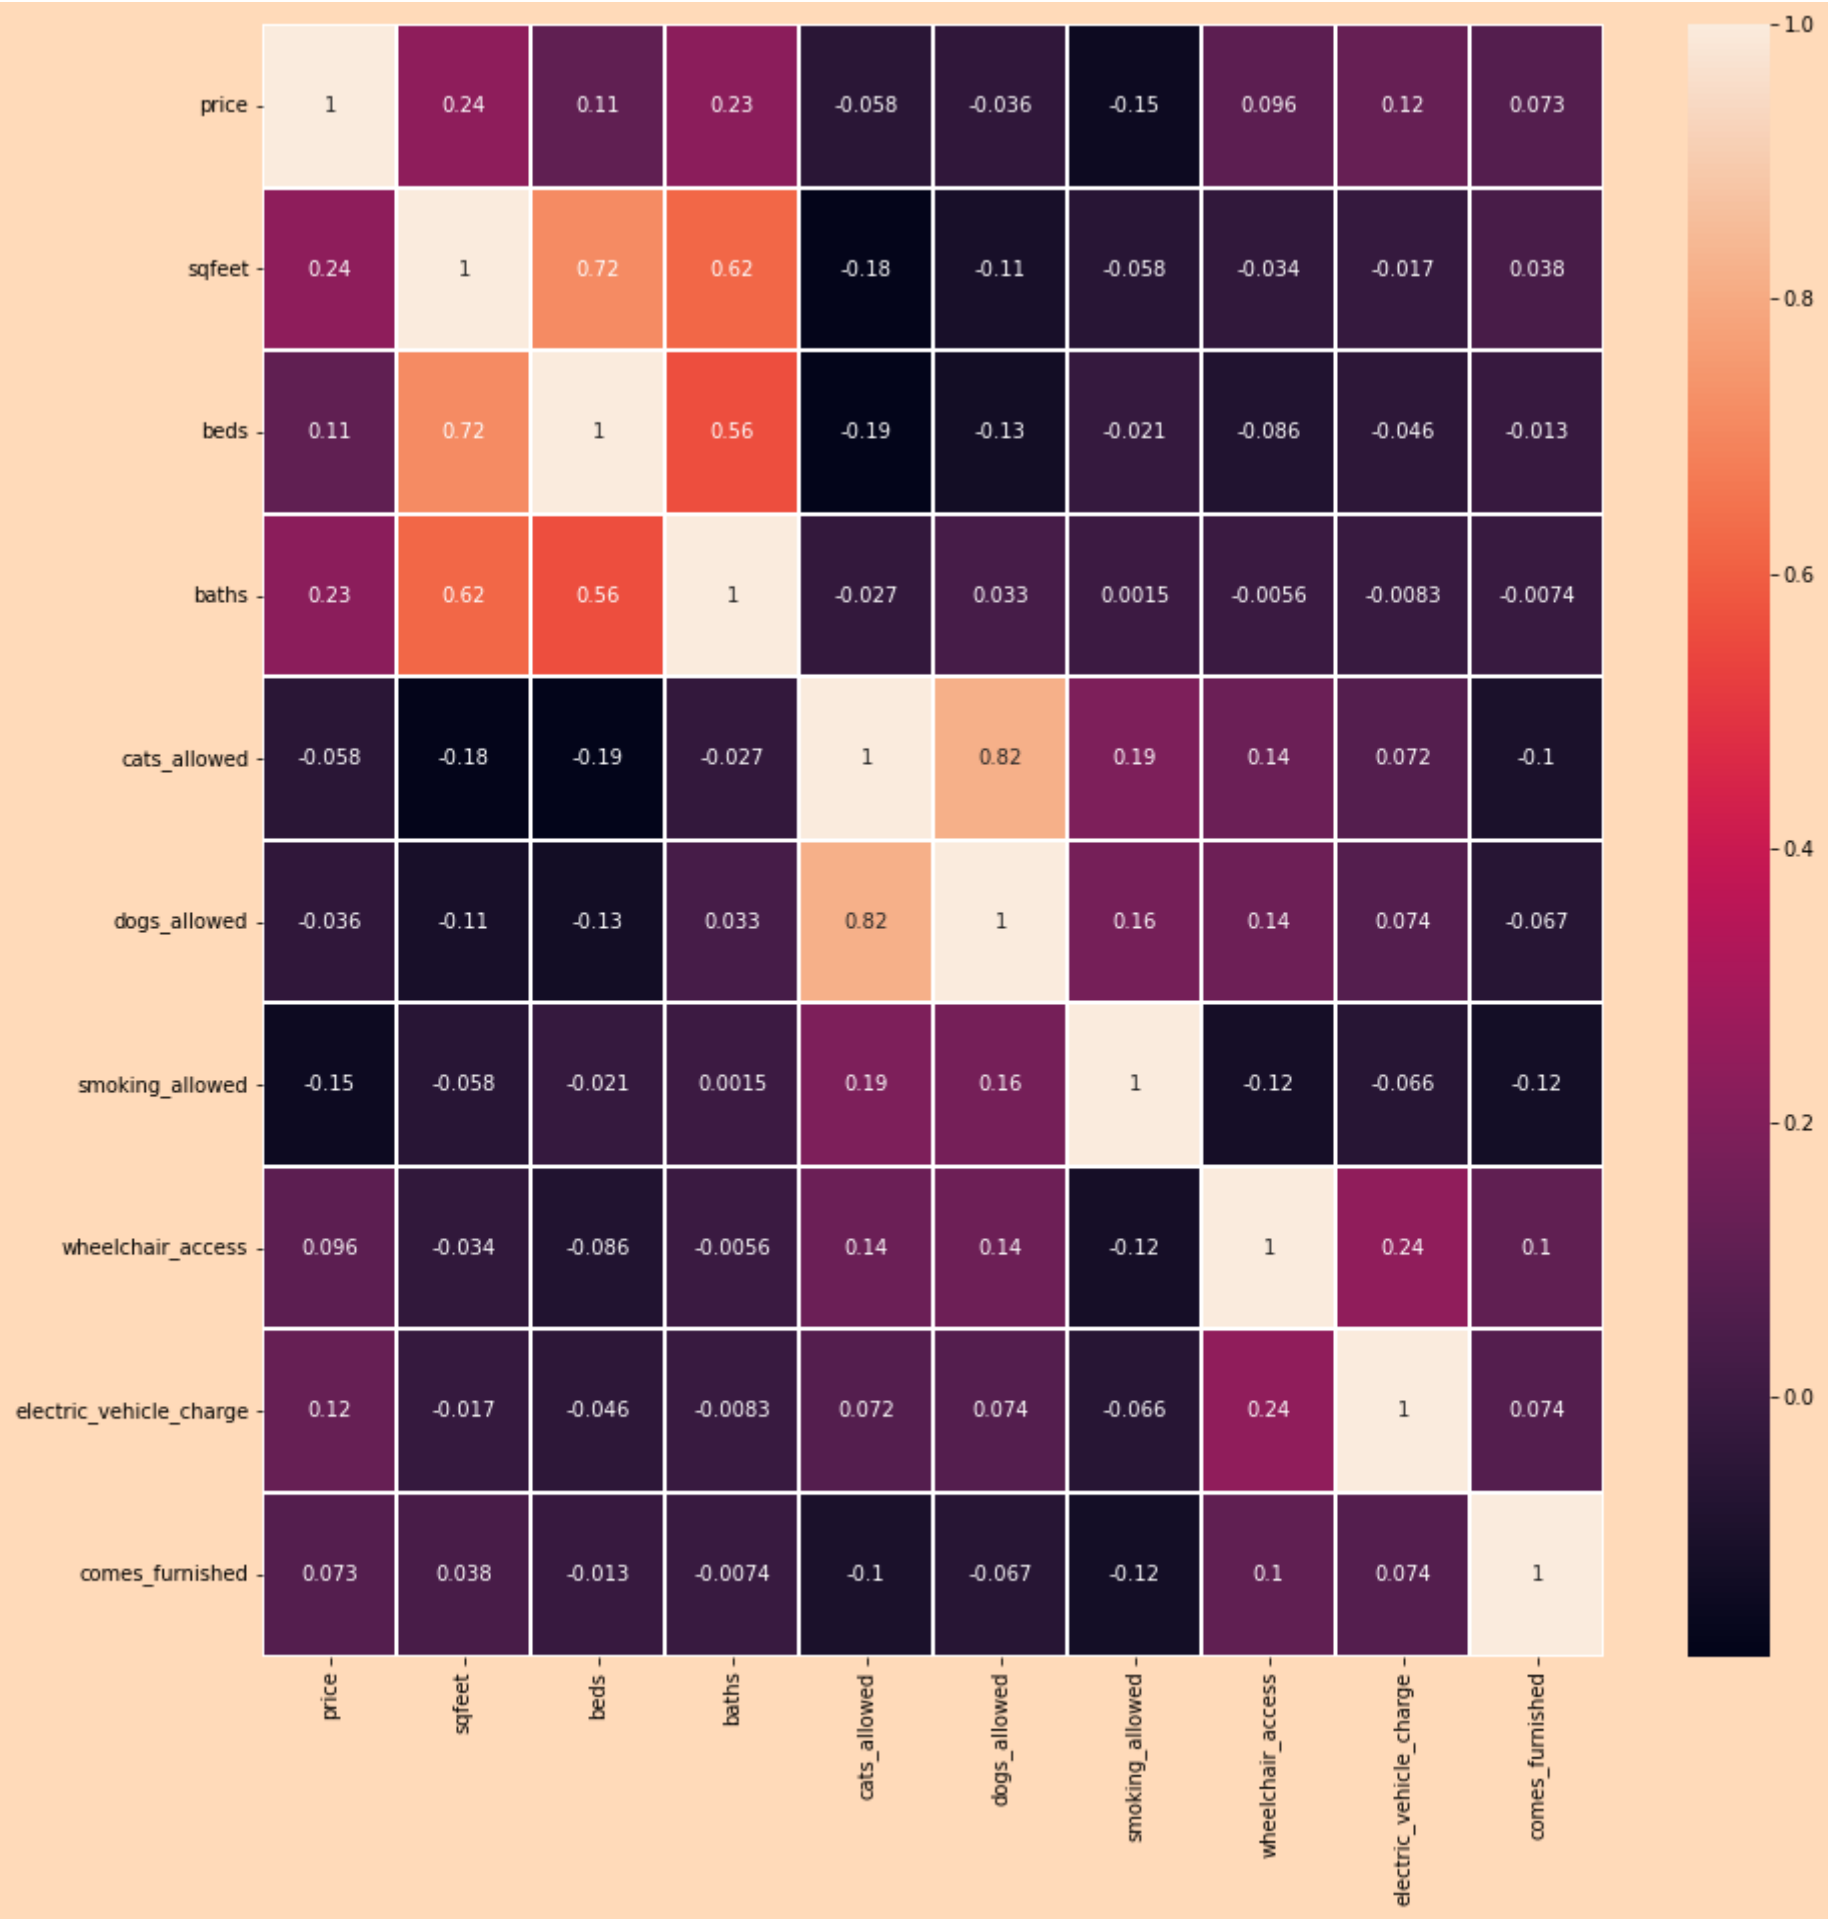
        plotnumber+=1
plt.tight_layout()
plt.show()
```



Using Heat map and Correlations

Lets check the dependency of columns with each others

```
In [15]: plt.figure(figsize=(15,15), facecolor='#ffdab9')
sns.heatmap(house.corr(), annot=True, linewidths=0.005);
```



Lets use get_dummies and convert object data to do further model application

```
In [16]: house = pd.get_dummies(data=house)
house.shape
```

Out[16]: (58974, 363)

linitializing Linear Model

```
In [17]: lr = LinearRegression()
```

Simple linear model

Lets check the result

```
In [18]: # Separating feature and output/result
X = house[['sqfeet']]
y = house[['price']]

# Splting train and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=20)

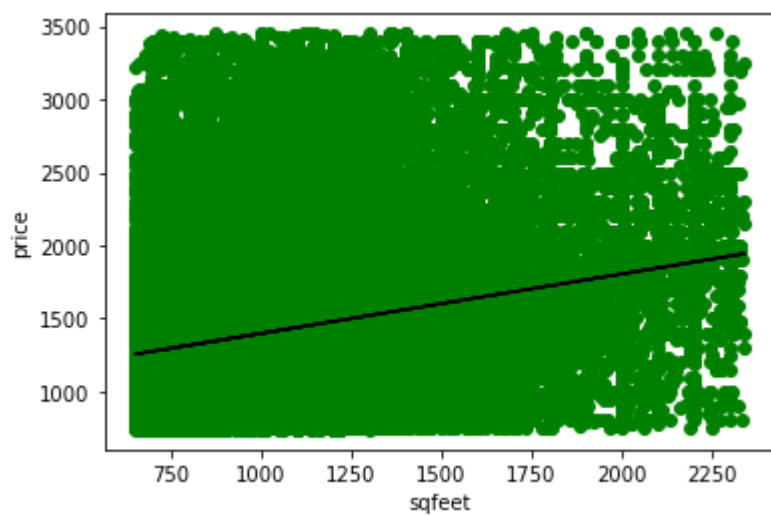
# Fitting linear modle
lr.fit(X_train, y_train)

print('Intercept',lr.intercept_)
print('Coefficient ',lr.coef_)
```

```
Intercept [991.48092574]
Coefficient [[0.40689659]]
```

Lets plot the scatter plot and see the predicted line

```
In [19]: plt.scatter(X, y, color='g')
plt.plot(X_test, lr.predict(X_test), color='k')
plt.xlabel('sqfeet')
plt.ylabel('price');
```



Lets see the MAE, MSE, RMSE and R2Score the view the model perfomance

```
In [20]: # Predict the model
y_pred= lr.predict(X_test)

print('Mean Absolute Error',mean_absolute_error(y_test, y_pred))
print('Mean Squared Error',mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error',np.sqrt(mean_squared_error(y_test, y_pred)))
print('R2 Score',r2_score(y_test, y_pred))
```

```
Mean Absolute Error 426.76385520718
Mean Squared Error 300213.8346956622
Root Mean Squared Error 547.9177262104797
R2 Score 0.055662803476350886
```

R2 score is very less, Lets do Multi Linear Regression and check R2 score

Multi Linear Regression

```
In [21]: # Sepearting Feature and result
X= house[['sqfeet','beds','baths']]
y = house[['price']]

# Spliting data into train and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=20)

# Applying linear model to data
lr.fit(X_train, y_train)

print('Intercept',lr.intercept_)
print('Coefficient ',lr.coef_)

# Predict the model
y_pred= lr.predict(X_test)

print('Mean Absolute Error',mean_absolute_error(y_test, y_pred))
print('Mean Squared Error',mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error',np.sqrt(mean_squared_error(y_test, y_pred)))
print('R2 score',r2_score(y_test, y_pred))
```

```
Intercept [961.3255184]
Coefficient [[ 0.42499237 -114.16746469 154.85637623]]
Mean Absolute Error 421.019551527462
Mean Squared Error 293008.7912622138
Root Mean Squared Error 541.3028646351447
R2 score 0.07832661749968284
```

R2 score has not improved much,

Lets make data scaler and apply Multi Regreesion model too

```
In [22]: # Initialize Scaler Model
scaler = StandardScaler()

# Apply scaler model
X_scaled = scaler.fit_transform(X)

# Splitting data into train and test data
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=200)

# Applying linear model to data
lr.fit(X_train, y_train)

print('Intercept',lr.intercept_)
print('Coefficient ',lr.coef_)

# Predict the model
y_pred= lr.predict(X_test)

print('Mean Absolute Error',mean_absolute_error(y_test, y_pred))
print('Mean Squared Error',mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error',np.sqrt(mean_squared_error(y_test, y_pred)))
print('R2 score',r2_score(y_test, y_pred))
```

```
Intercept [1437.43484598]
Coefficient [[142.98260846 -92.13740545  90.10055812]]
Mean Absolute Error 418.1886532532226
Mean Squared Error 288534.5851372588
Root Mean Squared Error 537.1541539793384
R2 score 0.08115875324956545
```

R2 score has not improved much,

Lets make data scaler and apply Decision Tree Regreesion model

Decision Tree

```
In [23]: # Sepearting Feature and result
X= house.drop(['price'], axis=1)
y = house[['price']]

# Apply scaler model
X_scaled = scaler.fit_transform(X)

# Splitting data into train and test data
X_train,X_test,y_train,y_test = train_test_split(X_scaled,y, test_size=0.2, random_state=100)

# Initializing Decision Tree
reg_tree = DecisionTreeRegressor(criterion= 'squared_error', max_depth= 3)

# Applying Decision Tree model to data
reg_tree.fit(X_train,y_train)

# Predict the model
y_pred = reg_tree.predict(X_test)

print('Mean Absolute Error',mean_absolute_error(y_test, y_pred))
print('Mean Squared Error',mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error',np.sqrt(mean_squared_error(y_test, y_pred)))
print('R2 score',r2_score(y_test, y_pred))
```

```
Mean Absolute Error 361.91216965499547
Mean Squared Error 224819.10073259717
Root Mean Squared Error 474.15092611171514
R2 score 0.31015164172673015
```

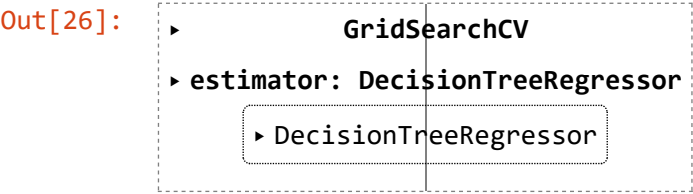
R2 score has improved very little, Lets do HyperParamete tuning and check best fit values to improve the R2

Hyperparameter Tuning

```
In [24]: grid_param = {
    'criterion': ['squared_error','friedman_mse'],
    'max_depth' : range(20,22),
    'min_samples_leaf' : range(12,18),
}
```

```
In [25]: grid_search = GridSearchCV(estimator=reg_tree,
    param_grid=grid_param,
    cv=5,
    n_jobs =-1)
```

```
In [26]: grid_search.fit(X_train,y_train)
```



```
In [27]: print(grid_search.best_params_)

{'criterion': 'squared_error', 'max_depth': 21, 'min_samples_leaf': 15}
```

```
In [28]: grid_search.best_score_
```

Out[28]: 0.6552040149302076

Applying Grid Search values to verify the result and check model score

```
In [29]: # Initializing Decision Tree
reg_tree = DecisionTreeRegressor(criterion= 'squared_error', max_depth= 21, min_samples_leaf=15 )

# Applying Decision Tree model to data
reg_tree.fit(X_train,y_train)

# Predict the model
y_pred = reg_tree.predict(X_test)

print('Mean Absolute Error',mean_absolute_error(y_test, y_pred))
print('Mean Squared Error',mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error',np.sqrt(mean_squared_error(y_test, y_pred)))
print('R2 score',r2_score(y_test, y_pred))
```

Mean Absolute Error 234.4015763883326
Mean Squared Error 110268.52863392427
Root Mean Squared Error 332.06705442413926
R2 score 0.6616454598410717

```
In [ ]:
```